

DATA ANALYSIS - SCIENTIFIC COMPUTING

Practical Tutorial 2 - Subspace Iteration Methods

Freshman Year, Computer Science Department



Issam Habibi
Younes Saoudi
Hamza Mouddene

2019-2020

Contents

1	Limitations of the Power Method	2
2	Extending the Power Method to Compute Dominant Eigenspace Vectors	4
2.1	A Basic Method to Compute a Dominant Eigenspace: <code>subspace_iter_v0</code>	4
2.2	Improved Version Using Raleigh-Ritz Projection: <code>subspace_iter_v1</code>	6
3	Towards an Efficient Solver: <code>subspace_iter_v2</code> & <code>subspace_iter_v3</code>	7
3.1	Block approach (<code>subspace_iter_v2</code>)	7
3.2	Deflation Method (<code>subspace_iter_v3</code>)	8
4	Numerical experiments	10
5	Fortran/Matlab Source Code Listings	12
5.1	Subspace Iteration Version 0	12
5.2	Edited Power Method (Matlab)	15
5.3	Subspace Iteration Version 1 (Matlab)	16
5.4	Subspace Iteration Version 2	19

Limitations of the Power Method

Question 01

Algorithm 1 Vector Power Method

Input: Matrix $A \in \mathbb{R}^{n \times n}$
Output: $(\lambda_1; v_1)$ eigenpair associated to the largest (in module) eigenvalue.
 $v \in \mathbb{R}^n$ given.
 $\beta = v^T \cdot A \cdot v$
repeat
 $y = A \cdot v$
 $v = y / \|y\|$
 $\beta_{old} = \beta$
 $\beta = v^T \cdot A \cdot v$
until $|\beta - \beta_{old}| / |\beta_{old}| < \epsilon$
 $\lambda_1 = \beta$ and $v_1 = v$

```

10 First eigenvalues
=====
Eigenvalue 1: 1.000 accuracy : 0.990E-15 number of iterations : 9890
Eigenvalue 2: 0.997 accuracy : 0.976E-15 number of iterations : 306
Eigenvalue 3: 0.903 accuracy : 0.977E-15 number of iterations : 2156
Eigenvalue 4: 0.891 accuracy : 0.982E-15 number of iterations : 1844
Eigenvalue 5: 0.878 accuracy : 0.912E-15 number of iterations : 275
Eigenvalue 6: 0.781 accuracy : 0.955E-15 number of iterations : 285
Eigenvalue 7: 0.702 accuracy : 0.931E-15 number of iterations : 493
Eigenvalue 8: 0.653 accuracy : 0.967E-15 number of iterations : 2304
Eigenvalue 9: 0.644 accuracy : 0.986E-15 number of iterations : 196
Eigenvalue 10: 0.553 accuracy : 0.974E-15 number of iterations : 538

=====
Time = 3.0499999523162842
=====

```

(a) Deflated Power Method

```

Calling LAPACK DSYEV on A
End of LAPACK DSYEV

10 First eigenvalues
=====
Eigenvalue 1: 1.000
Eigenvalue 2: 0.997
Eigenvalue 3: 0.903
Eigenvalue 4: 0.891
Eigenvalue 5: 0.878
Eigenvalue 6: 0.781
Eigenvalue 7: 0.702
Eigenvalue 8: 0.653
Eigenvalue 9: 0.644
Eigenvalue 10: 0.553

=====
Time = 0.10800000280141830
=====

```

(b) LAPACK subroutine `dsyev`

Figure 1.1: Comparison between the running times of the subroutine `deflated_power_method` and the Lapack subroutine `dsyev` when computing 10 eigenpairs.

One can see from the figures that, although the extracted eigenvalues are obviously the same, the Deflated Power Method is 30 times slower than the `dsyev` LAPACK subroutine when trying to compute 10 eigenpairs.

Question 02

The main drawback of the Deflated Power Method is the fact that this method computes the whole spectral decomposition of the matrix A , which is a square $n \times n$ matrix; a process which proves to be exceptionally time consuming. For the LAPACK subroutine `dsyev`, the matrix A is first reduced to tridiagonal form using

orthogonal similarity transformations, then the QR algorithm is applied to the tridiagonal matrix to compute the eigenvalues which is significantly more optimized.

Extending the Power Method to Compute Dominant Eigenspace Vectors

2.1 A Basic Method to Compute a Dominant Eigenspace: `subspace_iter_v0`

Question 03

If one were to extend Algorithm 1 to iterate on the matrix V , then the latter will **not** converge towards a matrix with columns containing the m most dominant eigenvectors as one can see in the following example:

Let there be a matrix $A = \begin{pmatrix} 7 & 9 & 9 \\ 9 & 7 & 9 \\ 9 & 9 & 7 \end{pmatrix}$. The two most dominant eigenpairs of the matrix A are : $\lambda_1 = 25$;

$$v_1 = \begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix} \text{ and } \lambda_2 = -2; v_2 = \begin{pmatrix} 0.4359 \\ -1.3831 \\ 0.9472 \end{pmatrix}$$

Trying to extract the two most dominant eigenvectors with the `edited_power_method.m` script will return

the dominant eigenvalue $\lambda_1 = 25$, which normally corresponds to the eigenvector $v_1 = \begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix}$, and the vectors

$$V = \begin{pmatrix} -0.5738 & -0.0636 \\ -0.5738 & -0.0636 \\ -0.5738 & -0.0636 \end{pmatrix}.$$

One can thus notice that the matrix V converges towards a matrix with columns just representing the eigenvector v_1 *under a different form* (which will be revealed shortly) instead of a matrix with columns containing the $m = 2$ most dominant eigenvectors.

This "different form" is actually none other than `beta/norm(beta)` with `beta = transpose(v) * A * v`.

If one computes its value after the while loop, they will find it to be equal to V .

To avoid this problem, `v = y/norm(y)` should be replaced with `v = mgs(y)` in the `edited_power_method.m` script (line 16). The matrix V will contain the eigenvectors associated to the m most dominant eigenvalues as a result. This, however, does not solve the drawback of the Power Method seen in **Question 01**.

(See Matlab Script `edited_power_method.m`)

Question 04

Although a computation of the whole spectral decomposition of the matrix H is performed in Algorithm 2, H is a square $m \times m$ matrix (with m being the number of eigenvalues extracted). Thus, this computation is far less significant, especially when compared to the one of the matrix A which is a square $n \times n$ matrix ($n > m$). The reason behind that is the fact that we only need a few number of eigenpairs in order to provide enough information about the data (as seen in TP1). Therefore, a data matrix containing thousands of columns could be analysed using only a few dozens of eigenpairs which leads us to choose a relatively small number m .

Question 05

Algorithm 2 Subspace Iteration Version 0

Input: Symmetric Matrix $A \in \mathbb{R}^{n \times n}$, the number of required eigenpairs m , the tolerance ϵ and the maximum number of iterations $MaxIter$.

Output: m dominant eigenvectors V_{out} and the corresponding eigenvalues Λ_{out}

Generate a set of m orthonormal vectors $V \in \mathbb{R}^{n \times m}$; $k = 0$

repeat

$k = k + 1$

$Y = A \cdot V$

$H = V^T \cdot A \cdot V$

Compute $acc = \|A \cdot V - V \cdot H\| / \|A\|$

$V \leftarrow$ orthonormalisation of the columns of Y

until ($k > MaxIter$ **or** $acc \leq \epsilon$)

Compute the spectral decomposition $X \cdot \Lambda_{out} \cdot X^T = H$, where the eigenvalues of H ($\text{diag}(\Lambda_{out})$) are arranged in descending order of magnitude.

Compute the corresponding eigenspace $V_{out} = V \cdot X$

(*See Fortran Script iter_v0.f90*)

2.2 Improved Version Using Raleigh-Ritz Projection: subspace_iter_v1

Question 06

Algorithm 3 Rayleigh-Ritz Projection

Input: Matrix $A \in \mathbb{R}^{n \times n}$ and an orthonormal set of Vectors V .
Output: The approximate eigenvectors V_{out} and the corresponding eigenvalues Λ_{out}
Compute the Rayleigh quotient $H = V^T \cdot A \cdot V$
Compute the spectral decomposition $X \cdot \Lambda_{out} \cdot X^T = H$, where the eigenvalues of H ($\text{diag}(\Lambda_{out})$) are arranged in descending order of magnitude.
Compute $V_{out} = V \cdot X$

Algorithm 4 Subspace Iteration Version 1 with Rayleigh-Ritz Projection

Input: Symmetric Matrix $A \in \mathbb{R}^{n \times n}$, the number of required eigenpairs m , the tolerance ϵ , the maximum number of iterations $MaxIter$ and the target percentage of the trace of A , $TracePercentage$
Output: n_{ev} dominant eigenvectors V_{out} and the corresponding eigenvalues Λ_{out}

Generate a set of m orthonormal vectors $V \in \mathbb{R}^{n \times m}$; $k = 0$; $PercentageReached = 0$
repeat
 $k = k + 1$
 Compute Y such that $Y = A \cdot V$
 $V \leftarrow$ orthonormalisation of the columns of Y
 Apply *Rayleigh-Ritz projection* on matrix A and orthonormal vectors V
 Convergence Analysis: save the converged eigenpairs and update $PercentageReached$
until ($PercentageReached > TracePercentage$ **or** $n_{ev} > m$ **or** $k > MaxIter$)

All of the steps of Algorithm 4 are already identified and highlighted in French in the file

subspace_iter_v1.m

(See *Matlab Script* subspace_iter_v1.m)

Towards an Efficient Solver: `subspace_iter_v2` & `subspace_iter_v3`

Two ways of improving the efficiency of the solver are proposed. Our aim is to build an algorithm that combines both the block approach and the deflation method in order to speed-up the convergence of the solver.

3.1 Block approach (`subspace_iter_v2`)

Orthonormality is performed at each iteration and is quite costly. One simple way to accelerate the approach is to perform p products at each iteration (replace $V = A.V$ (first step of the iteration) by $V = A^p.V$). Note that this very simple acceleration method is applicable to all versions of the algorithm.

Question 07

The definition of matrix multiplication is that if $C = AB$ for an $n \times m$ matrix A and an $m \times p$ matrix B , then C is an $n \times p$ matrix with entries

$$c_{ij} = \sum_{k=1}^m a_{ik}b_{kj} \quad (3.1)$$

From this, a simple algorithm can be constructed which loops over the indices i from 1 through n and j from 1 through p , computing the above using a nested loop:

Algorithm 5 Iterative algorithm of matrix multiplication

```
1: Input: matrices  $A$  and  $B$ 
2: Let  $C$  be a new matrix of the appropriate size
3: for  $i$  from 1 to  $n$  do
4:   for  $j$  from 1 to  $p$  do
5:      $sum = 0$ 
6:     for  $k$  from 1 to  $m$  do
7:        $sum \leftarrow sum + A_{ik} * B_{kj}$ 
8:     end for
9:      $C_{ij} \leftarrow sum$ 
10:  end for
11: end for
12: return  $C$ 
```

This algorithm takes time $\Theta(n \times m \times p)$. A common simplification for the purpose of algorithms analysis is to assume that the inputs are all square matrices of size $n \times n$, in which case the running time is $\Theta(n^3)$, i.e., cubic. So if we assume that A is equal to B we will find that the calculation A^2 has a complexity of $\Theta(n^3)$, then by recurrence we can show that the calculation A^p has a complexity of $\Theta((p-1)n^3)$.

Now we suppose that we have a matrix $V \in M_{n \times m}$, the calculation of $A^p.V$ has a complexity of $\Theta(m \times n^2 + (p-1)n^3)$.

One good way to reduce the cost is implementing a loop that calculates $V = AV$ p times rather than calculating A^p , this will lead us to the same result $V = A^p V$ in a less costly way due to the size of V , the complexity here should be $\Theta(p \times m \times n^2)$.

Another way out is Singular Value Decomposition SVD. Given an $n \times n$ real matrix A of full rank, SVD splits it apart as $A = U\Sigma U^T$ where Σ is a diagonal matrix, in time $\Theta(n^3)$. By the properties of SVD, $A^m = U\Sigma^m U^T$, so only the powers of the diagonal matrix need to be computed, and this can be done in $\Theta(n \log(m))$ time.

Performing the final multiplication $U \times \Sigma^m \times U^T$ takes $\Theta(n^{2.3727})$, so we have altogether $\Theta(n^3 + n \log(m))$ operations.

Question 08

iter_v2.f90
(See Fortran file iter_v2.f90)

3.2 Deflation Method (subspace_iter_v3)

Because the columns of V converge in order, we can freeze the converged columns of V . This freezing results in significant savings in the matrix-vector ($V = A\Delta V$), the orthonormalisation and Rayleigh-Ritz Projection steps.

Specifically, suppose the first nbc^2 columns of V have converged, and partition $V = [V_c, V_{nc}]$ where V_c has nbc columns and V_{nc} has $m - nbc$ columns. Then, we can form the matrix $[V_c, A\Delta V_{nc}]$, which is the same as if we multiply V by A . However, we still need to orthogonalise V_{nc} with respect to the frozen vectors V_c by first orthogonalising V_{nc} against V_c and then against itself.

Finally, the Rayleigh-Ritz Projection step can also be limited to the columns V_{nc} of V .

Question 09

The execution of the main program with the option `disp = 2`, displays the precision of a vector when it converges at the end of the computation to check if everything is OK. With the `subspace_iter_method_v1`, this precision is different for some of the vectors because of the Raleigh-Ritz projection which gives as output a vector containing the approximations of the eigenvalues which is used to calculate the norm of the residue, to determine if a vector converges or not, and thus to know if the eigenvalue makes it possible to reach the percentage.

Question 10

The subspace iter v2 method will be much faster than the subspace iter v1, because using $V = A^p$ instead of $V = A \times V$ allows us to accelerate the method globally, unlike the first one that uses an orthonormalization and computation of proper couples all the iterations.

Question 11

iter_v3.f90
(*See Fortran file iter_v3.f90*)

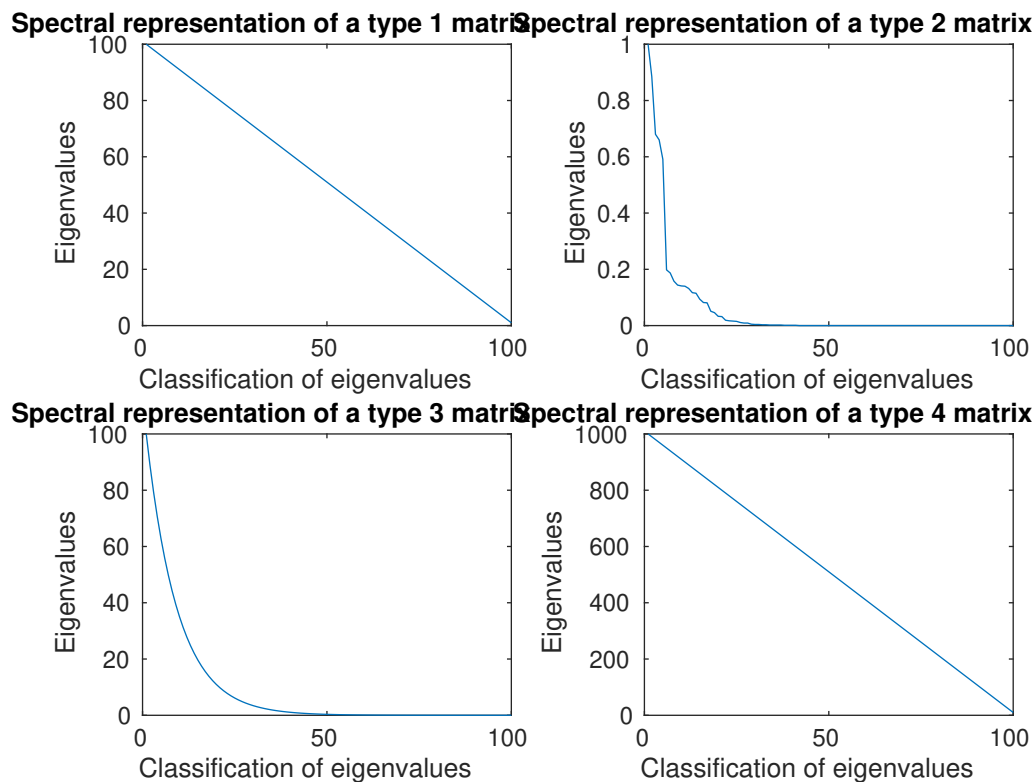
Numerical experiments

Question 12

We can notice that when p increases, time decreases as a logical consequence of the decrease in the number of iterations in our algorithm. As such, the reduction in orthonormalization and eigenvector search allows us to save considerable time.

Question 13

The following images show the eigenvalue distribution of these different types:



Question 14

For the first and the fourth types , the *DSYEV* is faster with a matrix that has big dimensions than the rest of the methods followed by *subspace_iter_v3* method . For matrix with small dimensions, all of the methods have almost the same performance with a small advantage again for the *DSYEV* method and the *power_v11* method. For the second and the third types , we can witness the same results except that *subspace_iter_v2* performs better in the second type.

Fortran/Matlab Source Code Listings

5.1 Subspace Iteration Version 0

```
1 ! This file is provided as part of the "projet long" for the "Calcul Scientifique et Analyse
   ! de Donn es" course
2 ! at INP-ENSEEIH
3 ! Date: 02/2019
4 ! Authors: P. Amestoy, P. Berger, A. Buttari, Y. Diouane, S. Gratton, R. Guivarch, F.H.
   ! Rouet, E. Simon
5 !
6 ! This file contains the implementation of
7 ! - one routine for computing the eigenpairs of a symmetric matrix with the basic subspace
   ! iteration method
8 ! (TO BE DEVELOPPED DURING "TP2 projet")
9 !
10 !
   ! -----
11 !
12 !! subspace_iter_v0
13 ! -----
14 !
15 ! routine that computes a certain number of eigenvalues and eigenvectors of a matrix A
16 ! using the subspace iteration method (v0)
17 !
18 ! ..Input arguments
19 !   n : integer, size of the matrix A
20 !   m : integer, maximum number of searched eigenvalues
21 !   a : double precision matrix
22 !
23 !   maxit : integer, maximum number of iterations of the power method
24 !   eps : the tolerance for the stopping criterion of the power method
25 !
26 ! ..Input/Output arguments
27 !   v : double precision matrix, (input) the starting subspace
28 !                                     (output) the eigenvectors corresponding
29 !                                     to the dominant eigenvalues
30 !
31 ! ..Output arguments
32 !   w : double precision vector, the eigenvalues
33 !   acc : double precision vector, accuracy for the invariant subspace
34 !   it : integer vector, number of iterations of the method
35 !
36 !   ierr : integer, indicates how the routine ends
37 !         0 : OK
38 !        -1 : inconsistency of the values of the arguments (ERROR)
39 !        -2 : allocation problem (ERROR)
40 !        -3 : maxit reached (ERROR)
41 !        -4 : problem in dsyev (ERROR)
42 !
```

```

-----
43 subroutine subspace_iter_v0(n, m, a, maxit, eps, v, w, acc, it, ierr)
44   implicit none
45   !! the subspace dimensions
46   integer,          intent(in)                :: n, m
47   !! the target matrix
48   double precision, dimension(n, n), intent(in) :: a
49   !! maximum # of iteration
50   integer,          intent(in)                :: maxit
51   !! the tolerance for the stopping criterion
52   double precision, intent(in)                :: eps
53   !! the m dominant eigenvalues
54   double precision, dimension(m), intent(out)   :: w
55   !! the starting subspace. The computed eigenvectors will be
56   !! returned in this array
57   double precision, dimension(n, m), intent(inout) :: v
58   !! the returned accuracy
59   double precision, intent(out)                :: acc
60   !! the number of iteration to converge
61   integer,          intent(out)                :: it
62   !! a flag for signaling errors
63   integer,          intent(out)                :: ierr
64
65   ! external functions
66   double precision, external                :: dlange
67
68   ! constants
69   integer, parameter                :: ione = 1
70   double precision, parameter       :: done = 1.d0, dzero = 0.d0, dmoins =
-1.d0
71
72   !! local variables
73   integer                :: i, j, k
74   double precision, allocatable, dimension(:, :) :: aux, y
75   !! h, rayleigh-Ritz quotient ; x, eigenvectors of h
76   double precision, allocatable, dimension(:, :) :: h, x
77   !! w_aux, eigenvalues computed by spectral decomposition subroutine DSYEV
78   double precision, allocatable, dimension(:) :: w_aux
79   double precision, allocatable, dimension(:) :: work
80   integer                :: lwork
81   double precision       :: normF_A
82
83
84   ierr = 0
85
86   if(m.gt.n)then
87     ierr = -1
88     return
89   end if
90
91   lwork = m*m + 5*m + n*n
92
93   allocate(aux(n, m), y(n, m), h(m, m), x(m, m), w_aux(m), work(lwork), stat = ierr)
94   if(ierr.ne. 0) then
95     ierr = -2
96     return
97   end if
98
99   normF_A = dlange('f', n, n, a, n, work)
100
101   k = 0
102   acc = 100*eps
103
104   !! v is a set of orthonormal vectors

```

```

105 call gram_schmidt(v, n, m, y)
106 v = y
107
108 do while((acc .ge. eps) .and. (k .lt. maxit))
109
110     k = k + 1
111
112     !! Compute y = a*v
113     call dgemm('n', 'n', n, m, n, done, a, n, v, n, dzero, y, n)
114
115     !! Compute h = v'*y
116     call dgemm('t', 'n', m, m, n, done, v, n, y, n, dzero, h, m)
117
118     !! Compute the accuracy ||a*v - v*h||/||A|| == ||y - v*h||/||A||
119     !! Compute aux = y-v*h
120     aux = y
121     call dgemm('n', 'n', n, m, m, dmoins, v, n, h, m, done, aux, n)
122
123     !! Compute acc = ||aux|| / ||A||
124     acc = dlange('f', n, m, aux, n, work) / normF_A
125
126     write(*, '(" IT:", i5, " -- Accuracy is: ", es10.2, a)', advance='no') k, acc, char(13)
127     !! V <- orthonormalization of y
128     call gram_schmidt(y, n, m, v)
129
130 end do
131
132 if(acc .lt. eps) then
133     ! compute the spectral decomposition of the Rayleigh quotient h
134     call dsyev('v', 'u', m, h, m, w_aux, work, lwork, ierr)
135     if( ierr .ne.0 )then
136         write(*, '("Error in dsyev")')
137         ierr = -4
138         goto 999
139     end if
140
141     !! Sort in the decreasing order (dsyev returns eigenvalues in ascending order)
142     !! (we suppose that all the eigen values are positive)
143     do i=1, m
144         x(:, i) = h(:, m-i +1)
145         w(i) = w_aux(m-i +1)
146     end do
147
148     !! v = v*x
149     y = v
150     call dgemm('n', 'n', n, m, m, done, y, n, x, m, dzero, v, n)
151
152 else
153     ierr = -3
154 end if
155
156 999 continue
157 deallocate(aux, y, h, x, w_aux, work)
158 write(*,*)
159
160 it = k
161
162 return
163 end subroutine subspace_iter_v0

```


5.2 Edited Power Method (Matlab)

```
1 function [ w, v0, it, flag ] = edited_power_method(A, m, eps, kmax)
2
3     n = size(A,1); %Number of row of Matrix A
4     %Question 3 Modification :
5     %OLD: v = ones(n,1); v = v / norm(v); % Random normalized Vector
6     %NEW:
7     v = randn(n, m); v = mgs(v); %Randomized orthonormal n*m Matrix
8
9
10    beta = norm(transpose(v)*A*v); %Eventual Eigenvalues
11    cv = false; %Did it converge?
12    k = 0; %Number of iterations
13
14    while(~cv)
15        y = A*v;
16        v = y/norm(y);
17        beta_old = beta;
18        %Question 3 Modification :
19        %OLD : beta = transpose(v) * y
20        %NEW:
21        beta = norm(transpose(v)*y);
22        err1 = (abs(beta - beta_old) / abs(beta_old));
23        k = k + 1;
24        cv = ( (err1 <= eps) | (k >= kmax));
25    end
26
27    it = k;
28    w = beta;
29    v0 = v;
30    if (cv)
31        flag = 0; %raise flag = 0 : converged!
32    else
33        flag = -3; %raise flag = -3 : could not converge!
34    end
```

5.3 Subspace Iteration Version 1 (Matlab)

```
1 % version améliorée de la méthode de l'espace invariant (v1)
2 % avec utilisation de la projection de Raleigh-Ritz
3
4 % Données
5 % A : matrice dont on cherche des couples propres
6 % m : taille maximale de l'espace invariant que l'on va utiliser
7 % percentage : pourcentage de la trace recherch
8 % eps : seuil pour déterminer si un vecteur de l'espace invariant a converg
9 % maxit : nombre maximum d'itération de la méthode
10
11 % Résultats
12 % W : vecteur contenant les valeurs propres (ordre d croissant)
13 % V : matrice des vecteurs propres correspondant
14 % n_ev : nombre de valeurs propres calculées
15 % it : nombre d'itérations de la méthode
16 % flag : indicateur sur la terminaison de l'algorithme
17 %   flag = 0 : on converge en ayant atteint le pourcentage de la trace recherch
18 %   flag = 1 : on converge en ayant atteint la taille maximale de l'espace
19 %   flag = -3 : on n'a pas converg en maxit itérations
20
21 function [ W, V, n_ev, it, flag ] = subspace_iter_v1( A, m, percentage, eps, maxit )
22
23     % calcul de la norme de A (pour le critère de convergence d'un vecteur (gamma))
24     normA = norm(A, 'fro');
25
26     % trace de A
27     traceA = trace(A);
28
29     % valeur correspondnat au pourcentage de la trace atteindre
30     vtrace = percentage*traceA;
31
32     n = size(A,1);
33     W = zeros(n,1);
34
35     % numéro de l'itération courante
36     k = 0;
37     % somme courante des valeurs propres
38     eigsum = 0.0;
39     % nombre de vecteurs ayant convergés
40     nb_c = 0;
41
42     % indicateur de la convergence
43     conv = 0;
44
45     % on génère un ensemble initial de m vecteurs orthogonaux
46     Vr = randn(n, m);
47     Vr = mgs(Vr);
48
49     % rappel : conv = (eigsum >= trace) | (nb_c == m)
50     while (~conv & k < maxit),
51
52         k = k+1;
53         %% Y <- A*V
54         Y = A*Vr;
55         %% orthogonalisation
56         Vr = mgs(Y);
57
58         %% Projection de Rayleigh-Ritz
59         [Wr, Vr] = rayleigh_ritz_projection(A, Vr);
60
61         %% Quels vecteurs ont converg cette itération
62         analyse_cvg_finie = 0;
```

```

63     % nombre de vecteurs ayant converg      cette it ration
64     nbc_k = 0;
65     % nb_c est le dernier vecteur      avoir converg      l'it ration pr c dente
66     i = nb_c + 1;
67
68     while(~analyse_cvg_finie),
69         % tous les vecteurs de notre sous-espace ont converg
70         % on a fini (sans avoir obtenu le pourcentage)
71         if(i > m)
72             analyse_cvg_finie = 1;
73         else
74             % est-ce que le vecteur i a converg
75
76             % calcul de la norme du r sidu
77             aux = A*Vr(:,i) - Wr(i)*Vr(:,i);
78             res = sqrt(aux'*aux);
79
80             if(res >= eps*normA)
81                 % le vecteur i n'a pas converg ,
82                 % on sait que les vecteurs suivants n'auront pas converg non plus
83                 % => it ration finie
84                 analyse_cvg_finie = 1;
85             else
86                 % le_vecteur i a converg
87                 % un de plus
88                 nbc_k = nbc_k + 1;
89                 % on le stocke ainsi que sa valeur propre
90                 W(i) = Wr(i);
91
92                 % on met      jour la somme des valeurs propres
93                 eigsum = eigsum + W(i);
94
95                 % si cette valeur propre permet d'atteindre le pourcentage
96                 % on a fini
97                 if(eigsum >= vtrace)
98                     analyse_cvg_finie = 1;
99                 else
100                     % on passe au vecteur suivant
101                     i = i + 1;
102                 end
103             end
104         end
105     end
106
107     % on met      jour le nombre de vecteurs ayant converg s
108     nb_c = nb_c + nbc_k;
109
110     % on a converg      dans l'un de ces deux cas
111     conv = (nb_c == m) | (eigsum >= vtrace);
112
113 end
114
115 if(conv)
116     % mise      jour des r sultats
117     n_ev = nb_c;
118     V = Vr(:, 1:n_ev);
119     W = W(1:n_ev);
120     it = k;
121 else
122     % on n'a pas converg
123     W = zeros(1,1);
124     V = zeros(1,1);
125     n_ev = 0;
126     it = k;
127 end

```

```
128
129     % on indique comment on a fini
130     if(eigsum >= vtrace)
131         flag = 0;
132     else if (n_ev == m)
133         flag = 1;
134     else
135         flag = -3;
136     end
137 end
138 end
```

5.4 Subspace Iteration Version 2

```

1 ! This file is provided as part of the "projet long" for the "Calcul Scientifique et Analyse
  ! de Donn es" course
2 ! at INP-ENSEEIH
3 ! Date: 02/2019
4 ! Authors: P. Amestoy, P. Berger, A. Buttari, Y. Diouane, S. Gratton, R. Guivarch, F.H.
  ! Rouet, E. Simon
5 !
6 ! This file contains the implementation of
7 ! - one routine for computing the eigenpairs of a symmetric matrix with the subspace
  ! iteration method
8 ! with Rayleigh-Ritz projection + block approach (v2) + deflation (v3) (TO BE DEVELOPPED)
9 !
10 !
  !-----
11 !
12 !! subspace_iter_v2
13 ! -----
14 !
15 ! routine that computes a certain amount of eigenvalues and eigenvectors of a matrix A
16 ! using the subspace iteration method with Rayleigh-Ritz projection + block approach (v2) +
  ! deflation (v3)
17 !
18 ! ..Input arguments
19 !   n : integer, size of the matrix A
20 !   m : integer, size of the searched subspace
21 !   a : double precision matrix
22 !   p : number of products  $Y = A^p \cdot V$ 
23 !
24 !   percentage : double precision, fraction of the trace we would like to obtain
25 !
26 !   maxit : integer, maximum number of iterations of the power method
27 !   eps : the tolerance for the stopping criterion of the power method
28 !
29 ! ..Input/Output arguments
30 !   v : double precision matrix, (input) the starting subspace
31 !                                     (output) the eigenvectors corresponding
32 !                                     to the dominant eigenvalues
33 !
34 ! ..Output arguments
35 !   w : double precision vector, the eigenvalues
36 !   n_ev : integer, the number of computed eigenvalues
37 !   acc_ev : double precision vector, accuracy for each eigenvalues
38 !   it_ev : integer vector, number of iterations for each eigenvalues
39 !
40 !   ierr : integer, indicates how the routine ends
41 !           0 : OK
42 !           1 : the maximum number of eigenvalues is reached before obtaining the percentage
  ! (WARNING)
43 !           -1 : inconsistency of the values of the arguments (ERROR)
44 !           -2 : allocation problem (ERROR)
45 !           -3 : maxit reached (ERROR)
46 !           -4 : problem in dsyev (ERROR)
47 !
  !-----
48 subroutine subspace_iter_v2(n, m, a, p, percentage, maxit, eps, v, w, n_ev, acc_ev, it_ev,
  ierr)
49 implicit none
50 !! the subspace dimensions
51 integer, intent(in) :: n, m
52 !! the target matrix

```

```

53 double precision, dimension(n, n), intent(in)      :: a
54 !! number of products  $Y = A^p \cdot V$ 
55 integer, intent(in)                               :: p
56 !! the percentage of the trace we want to obtain
57 double precision, intent(in)                       :: percentage
58 !! maximum # of iteration
59 integer, intent(in)                               :: maxit
60 !! the tolerance for the stopping criterion
61 double precision, intent(in)                       :: eps
62 !! the starting subspace. The computed eigenvectors will be
63 !! returned in this array
64 double precision, dimension(n, m), intent(inout) :: v
65 !! the m dominant eigenvalues
66 double precision, dimension(m), intent(out)        :: w
67 !! number of computed eigenvalues
68 integer, intent(out)                              :: n_ev
69 !! accuracy for each eigenvalues
70 double precision, dimension(m), intent(out)        :: acc_ev
71 !! number of iterations for each eigenvalues
72 integer, dimension(m), intent(out)                 :: it_ev
73 !! a flag for signaling errors
74 integer, intent(out)                              :: ierr
75
76 ! external functions
77 double precision, external                         :: dlange, ddot
78
79 ! constants
80 integer, parameter                                :: ione = 1
81 double precision, parameter                       :: done = 1.d0, dzero = 0.d0, dmoins =
-1.d0
82
83 !! local variables
84 integer                                           :: i, j
85 double precision, allocatable, dimension(:, :)   :: y
86 double precision, allocatable, dimension(:, :)   :: h, x
87 double precision, allocatable, dimension(:)       :: aux_acc, w_aux, t
88 double precision                                :: trace, p_trace, eig_sum, normF_A
89 integer                                           :: conv
90 integer                                           :: k
91 double precision                                :: acc
92 logical                                           :: ok
93 double precision                                :: beta
94 !! the length of the workspace for dsyev
95 integer                                           :: lwork
96 !! the workspace
97 double precision, allocatable, dimension(:)       :: work
98 !! the matrix where we will store  $A^p$ 
99 double precision, dimension(n,n)                 :: c
100
101 ierr = 0
102
103 !inconsistency of the values of the arguments (ERROR)
104 if(m.gt.n)then
105     ierr = -1
106     return
107 end if
108
109 !allocation problem (ERROR)
110 lwork = m*m + 5*m + n*n
111 allocate(y(n, m), h(m, m), x(m, m), w_aux(m), aux_acc(n), t(m), work(lwork), stat=ierr)
112 if(ierr .ne. 0) then
113     ierr = -2
114     return
115 end if
116

```

```

117  !! Calculate  $A^p$ 
118  c = a
119  do i = 1, p
120      call dgemm('n', 'n', n, n, n, done, a, n, c, n, dzero, c, n)
121  end do
122
123  !Calculate the trace of matrix A
124  trace = 0.d0
125  do i = 1, n
126      trace = trace + a(i,i)
127  end do
128  p_trace = percentage * trace
129  eig_sum = 0.d0
130
131  normF_A = dlange('f', n, n, a, n, work)
132
133  it_ev = 0
134  acc_ev = 0.D0
135  n_ev = 0
136  k = 0
137
138  !! Initial set of orthonormal vectors
139  call gram_schmidt(v, n, m, y)
140  v = y
141
142  do while((eig_sum .lt. p_trace) .and. (n_ev .lt. m) .and. (k .lt. maxit))
143
144      k = k + 1
145
146      !! A. Compute  $y = a*v$ 
147      call dgemm('n', 'n', n, m, n, done, c, n, v, n, dzero, y, n)
148
149      !! B. Orthonormalisation
150      call gram_schmidt(y, n, m, v)
151
152      !! C. Rayleigh-Ritz projection
153      !! 1.  $H = V^T A V$ 
154      !!  $Y = A V$ 
155      call dgemm('n', 'n', n, m, n, done, c, n, v, n, dzero, y, n)
156      !!  $H = V^T * Y$ 
157      call dgemm('t', 'n', m, m, n, done, v, n, y, n, dzero, h, m)
158      !! 2. Spectral decomposition
159      call dsyev('v', 'u', m, h, m, w_aux, work, lwork, ierr)
160      !problem in dsyev (ERROR)
161      if( ierr .ne.0 )then
162          write(*, '("Error in dsyev")')
163          ierr = -4
164          goto 999
165      end if
166
167      !! Sort in the decreasing order
168      !! (we suppose that all the eigen values are positive)
169      do i = 1, m
170          t(i) = w_aux(m-i+1)
171          x(:, i) = h(:, m-i+1)
172      end do
173
174      !! 3.  $V = VX$ 
175      y = v
176      call dgemm('n', 'n', n, m, m, done, y, n, x, m, dzero, v, n)
177
178      !! D. Convergence analysis step
179      conv = 0
180      i = n_ev + 1
181      !! the larger eigenvalue will converge more swiftly than

```

```

182      !! those corresponding to the smaller eigenvalue.
183      !! for this reason, we test the convergence in the order
184      !! i=1,2,.. and stop with the first one to fail the test
185      ok = .false.
186      do while(.not. ok)
187          if( i .gt. m) then
188              ok = .true.
189          else
190              !!compute acc=norm(a*v(:,i) - v(:,i)*t(i),2)/lambda;
191              !!--compute aux_acc=a*v(:,i) - v(:,i)*t(i)
192              !!--compute acc=||aux_acc||/||a||
193              aux_acc = v(:,i)
194              beta = -t(i)
195              call dgemv('n', n, n, done, a, n, v(1,i), ione, beta, aux_acc, ione)
196              acc = sqrt(ddot(n, aux_acc, ione, aux_acc, ione))/normF_A
197              ! write(*,*) i, acc
198
199              if(acc.gt.eps) then
200                  ok = .true.
201              else
202                  ! write(*,*) 'vector', i, 'converges', acc
203                  conv = conv + 1
204                  w(i) = t(i)
205                  acc_ev(i) = acc
206                  it_ev(i) = k
207                  eig_sum = eig_sum + w(i)
208                  i = i + 1
209                  if( eig_sum .ge. p_trace) ok = .true.
210              end if
211          end if
212      end do
213
214      n_ev = n_ev + conv
215      !write(*,*) n_ev
216
217  end do
218
219      !the maximum number of eigenvalues is reached before obtaining the percentage (WARNING)
220      if(n_ev .eq. m) then
221          ierr = 1
222      end if
223
224      !maxit reached (ERROR)
225      if(k .eq. maxit) then
226          ierr = -3
227      end if
228
229 999 continue
230      deallocate(y, h, x, w_aux, aux_acc, t)
231      write(*,*)
232
233      return
234  end subroutine subspace_iter_v2

```